ORDINA
a Sopra Steria company


# INFORMATION RETRIEVAL

## Techathon Data Science & AI

25/01/2024

Ordina-Group/hackathon-vectorsearch-benchmark: techathon Data Science & AI about vector search, embedding models and benchmarks (github.com)

By Pauline van Nies

# INFORMATION RETRIEVAL

Finding documents of text that satisfies the information need from within a large collection.

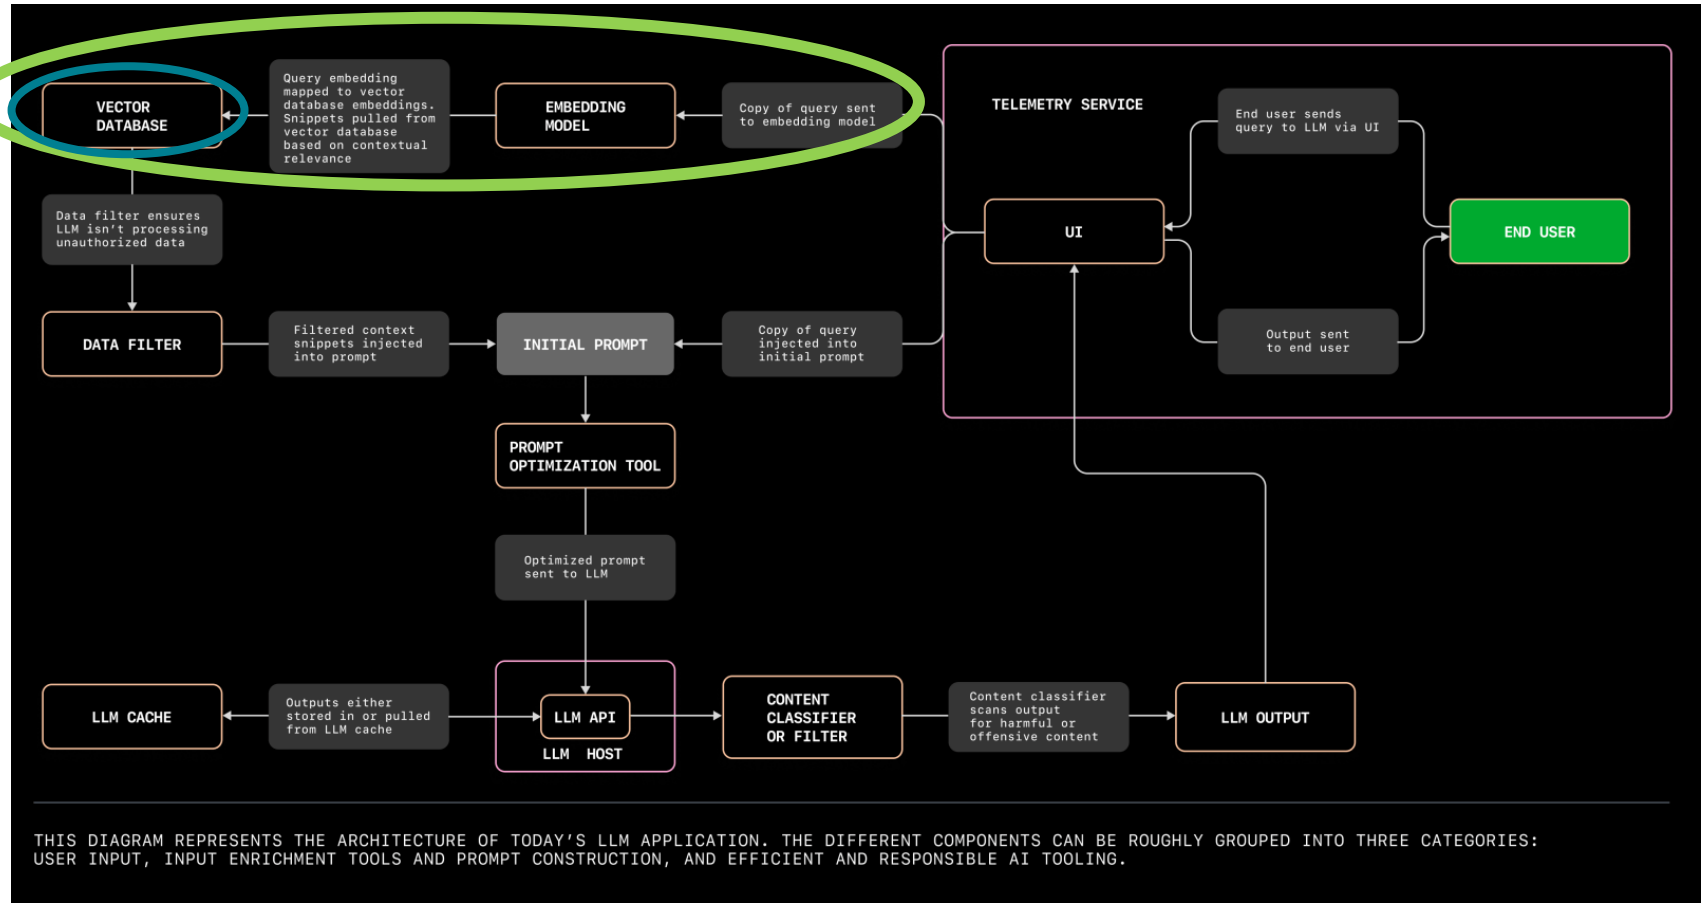


Bigger picture : The architecture of today's LLM applications - The GitHub Blog

# GOAL OF THIS TECHATHON

Have a good understanding of the components and concepts


Embedding models
Data to embed

Information retrieval
benchmarks

Vector database
Search settings

How?

- Working in groups investigating questions and being satisfied with the answers
- Check repo and run evaluation on benchmark dataset with model in weaviate

ORDINA
a Sopra Steria company

# EMBEDDING MODEL

## Text -> vector representation

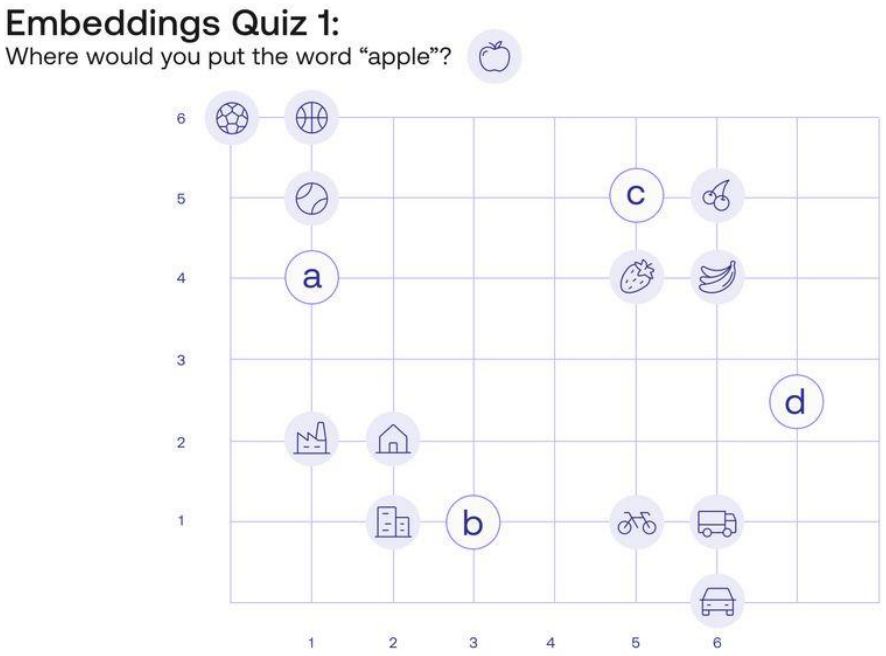


### Questions

- What is an embedding model?
- How is an embedding model trained?
- On what properties do the embedding models differ?
- Which embedding model would you like to test and why?

### Resources

The internet

Massive Text Embedding Benchmark Leaderboard

ORDINA
a Sopra Steria company

# COMPARISON OF EMBEDDING MODELS

## Tasks, datasets, evaluation metrics , BEIR and MTEB


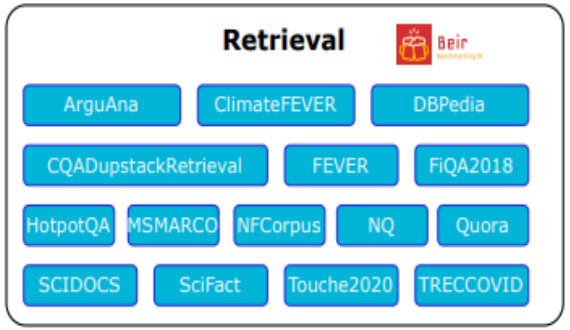


### Questions

- On which tasks are embedding models compared?
- What is the difference between reranking and retrieval tasks?
- What are the different datasets used for the retrieval tasks, its different properties?
- What is the format of the datasets? Understand corpus, queries, qrels.
- What are the evaluation metrics used for information retrieval in BEIR and in MTEB?
- Is there anything you find in the evaluation approach that stands out?
- Which dataset is interesting for your use case and why?

### Resources

Massive Text Embedding Benchmark Leaderboard
BEIR colab for exploration and evaluation
BEIR Paper
MTEB Paper

ORDINA
a Sopra Steria company

# VECTOR DATABASE

## Weaviate

### Questions

- What are alternatives to Weaviate and what are pros and cons?
- What does Weaviate embed in the vector?
- What can be changed in preprocessing and for different languages?
- What does hybrid search do?
- Which parameters are adjustable for hybrid search?
- What is the default distance metric for vector search?
- What are different ranking methods and how do they influence the vector search?
- Which parameters should be chosen in weaviate to reproduce a benchmark evaluation?

### Resources

- Weaviate documentation
  - Hybrid search
  - text2vec-transformers
  - Collection schema
  - Getting started with Weaviate Python Library

# PERFORMING EVALUATION

## Techathon repo

hackathon-vectorsearch-benchmark

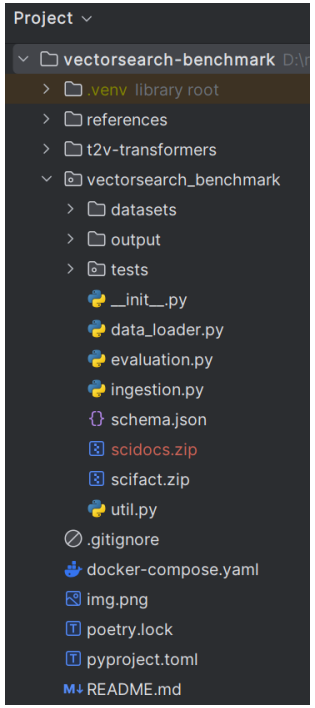


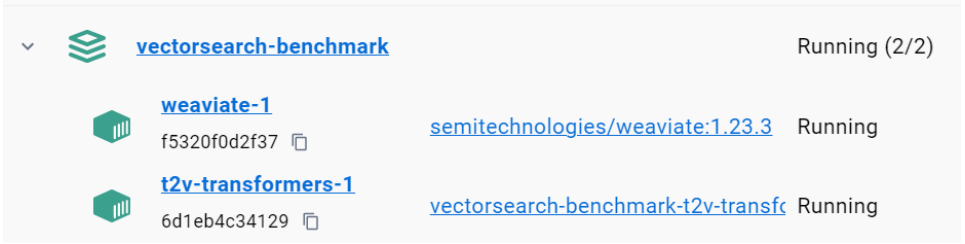


You can run an evaluation using this mini application consisting of the weaviate vector database filled with a benchmark dataset.

- unzip `scifact.zip` to a folder in `vectorsearch_benchmark` called `datasets` .
- perform the prerequisites as described above
- run `ingestion.py` e.g. by `poetry run python ingestion.py` in the `vectorsearch_benchmark` folder or using your IDE run configurations
- run `evaluation.py`

Note: Runs fine on my laptop but if you run into time out errors, try to reduce the `BATCH_SIZE` in `ingestion.py`

### Ideas

- Experiment with weaviate vector search parameters in `evaluation.py`
- Change the embedding model in the build args in the `docker-compose.yaml`
- Investigate another dataset

## Questions

- How does the evaluation metric compare to the benchmarks?
- If there is a discrepancy, do you have an idea where it might come from?

ORDINA
a Sopra Steria company